The power of random

By RandomBaseline

Ideas

- Dynamical Systems  ➔ not achievable in a weekend
- Road prediction + linear approximation  ➔ good road prediction not feasible
- Backpropagate the target function directly  ➔ no differentiable objective function
- ANN trained with Genetic Algorithm 

Genetic Algorithm – Advantages

- non-gradient based (T_{Target} is not differentiable)
- parallel search (less likely to get stuck in local optima)
- can optimize for multiple objectives (constraint & target)
- explainable method

Design – Code

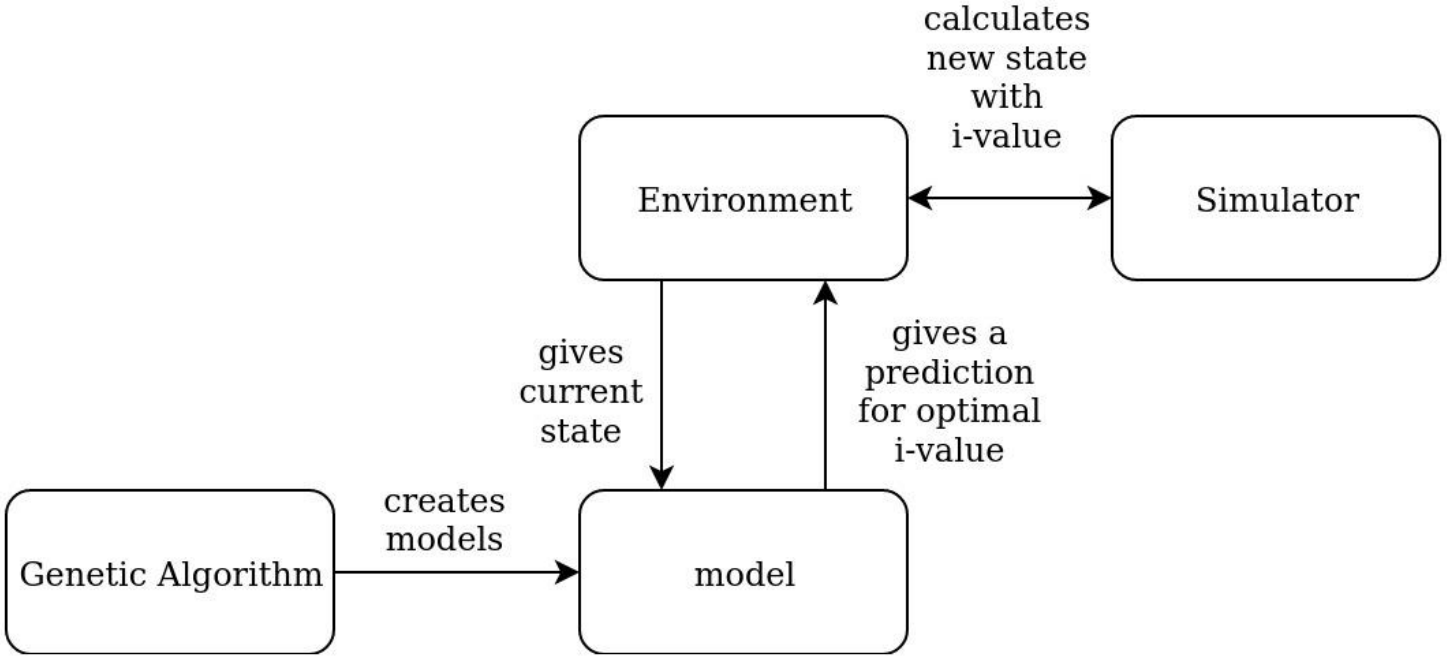
calculates
new state
with
i-value
Environment
Simulator
gives
current
state
gives a
prediction
for optimal
i-value
creates
models
Genetic Algorithm
model

Design – Artificial Neural Network (ANN)

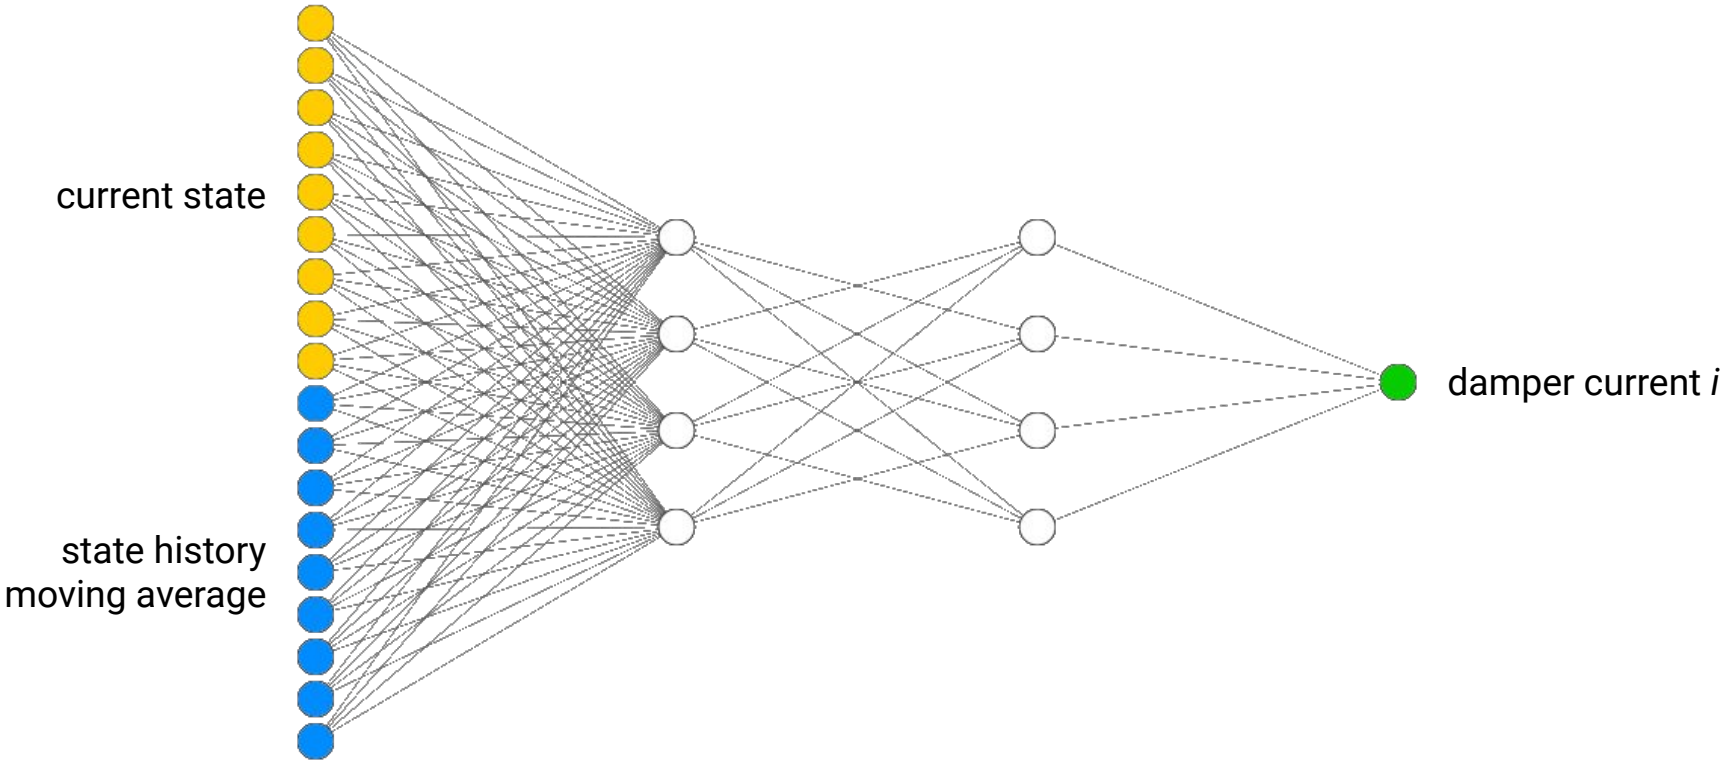

Design - Genetic Algorithm

Fitness function:

$$f(x) = \begin{cases} T_{Target} & constraint = true \\ T_{Target} \cdot 10 & else \end{cases}$$

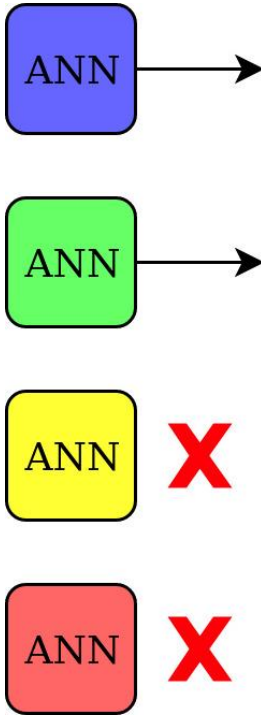

Design - Genetic Algorithm

Fitness function:

$$f(x) = \begin{cases} T_{Target} & constraint = true \\ T_{Target} \cdot 10 & else \end{cases}$$

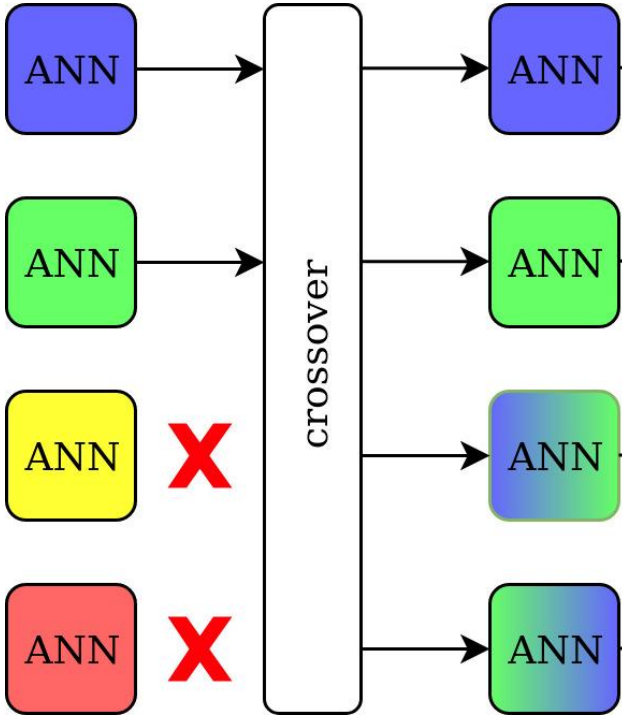

Design - Genetic Algorithm

Fitness function:

$$f(x) = \begin{cases} T_{Target} & constraint = true \\ T_{Target} \cdot 10 & else \end{cases}$$

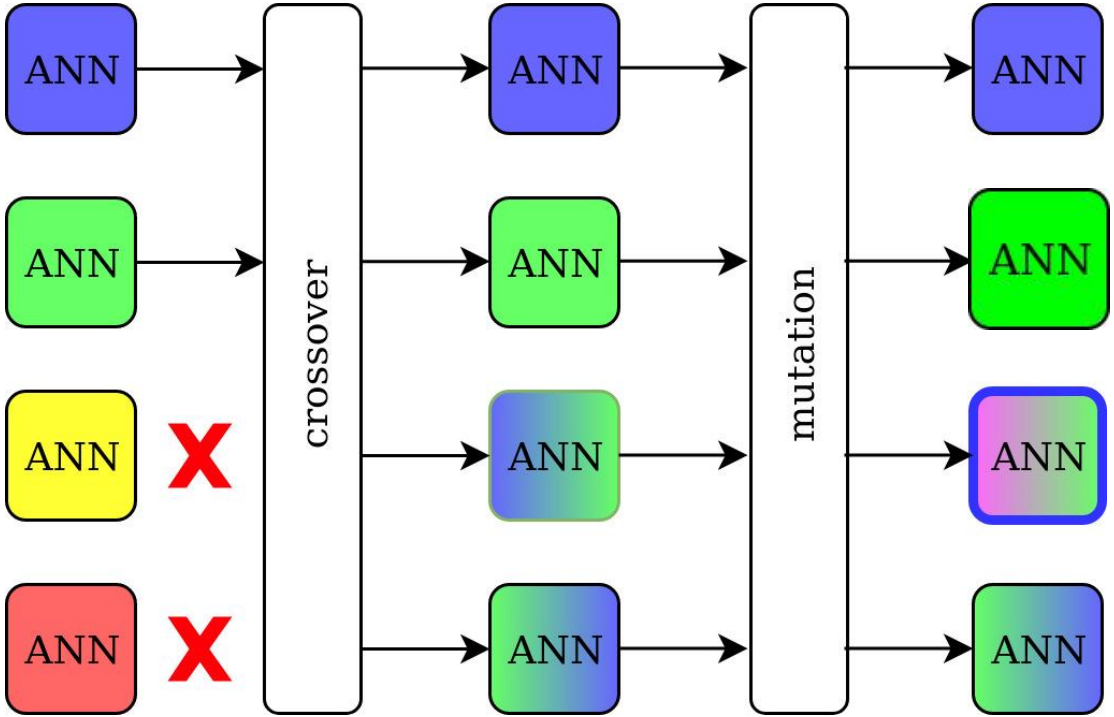

Training - Hyperparameters

- **EPOCHS:** in an epoch a population is evaluated and updated
- **EVALUATION REPEATS:** how many road positions we look at for each epoch
- **EVALUATION STEPS:** how many steps we train from the picked road position onward

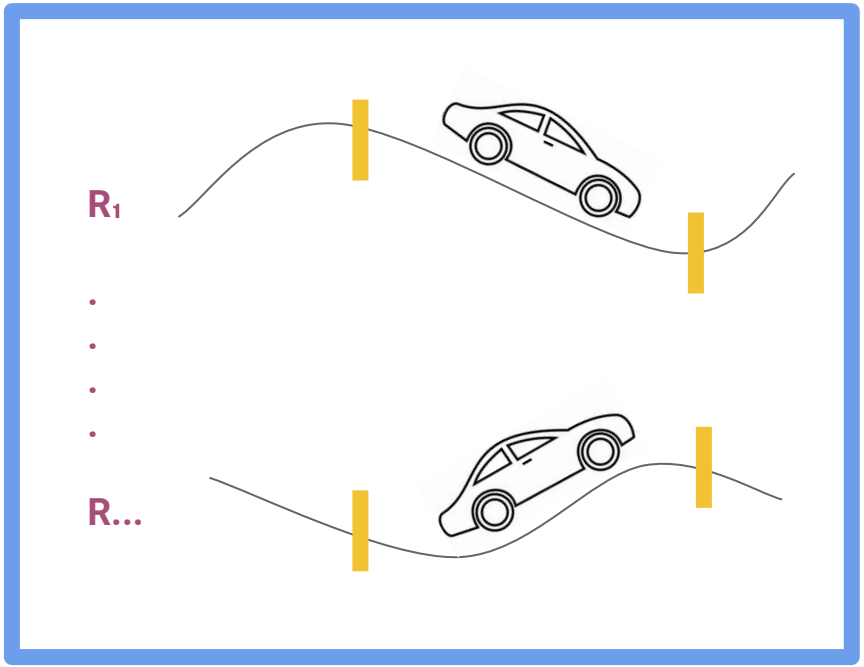

Results

- including a moving average improved performance

- adding convolution did not improve performance directly
  - might achieve better results after longer training

- our model achieves T_{Target} values lower than simulated passive dampers (constant *i*)

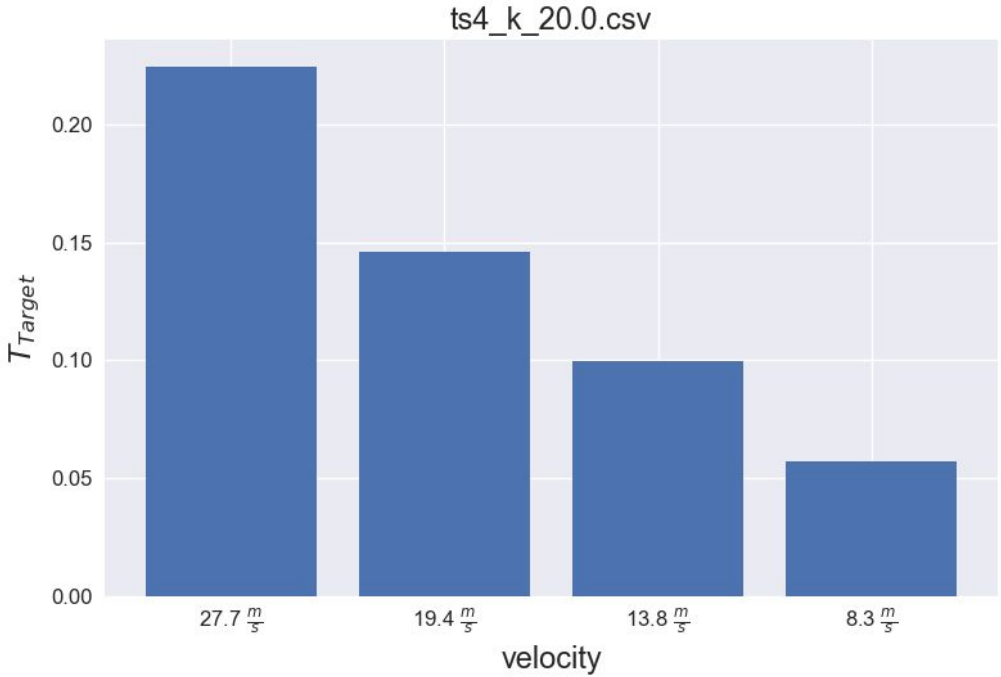

Constant vs adaptive i - Example

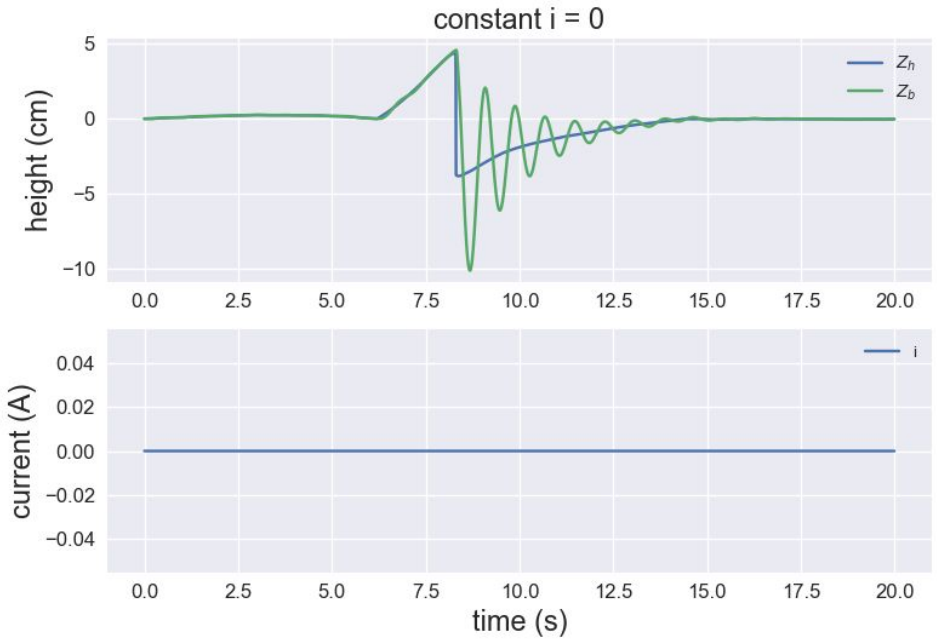

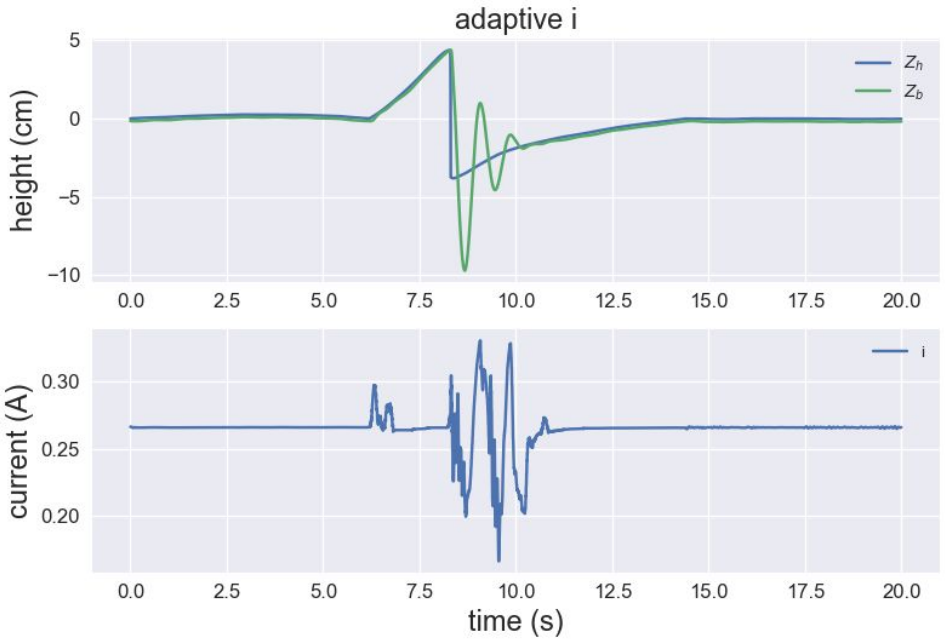

ANN inference speed

0,25ms

Outlook – Reinforcement Learning

- same environment can be used
- more directed learning → might converge faster
- challenge: continuous action space, delayed rewards

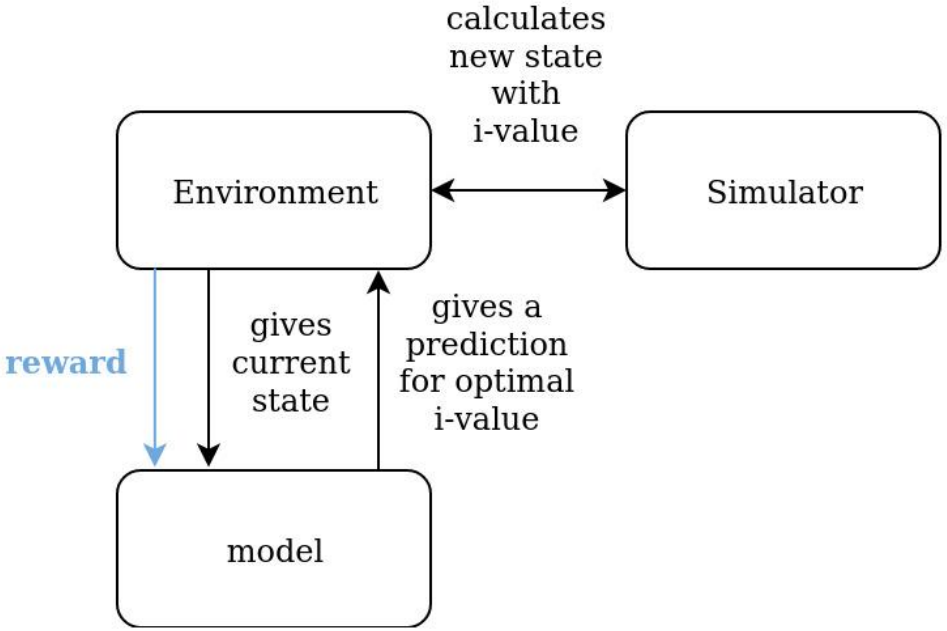

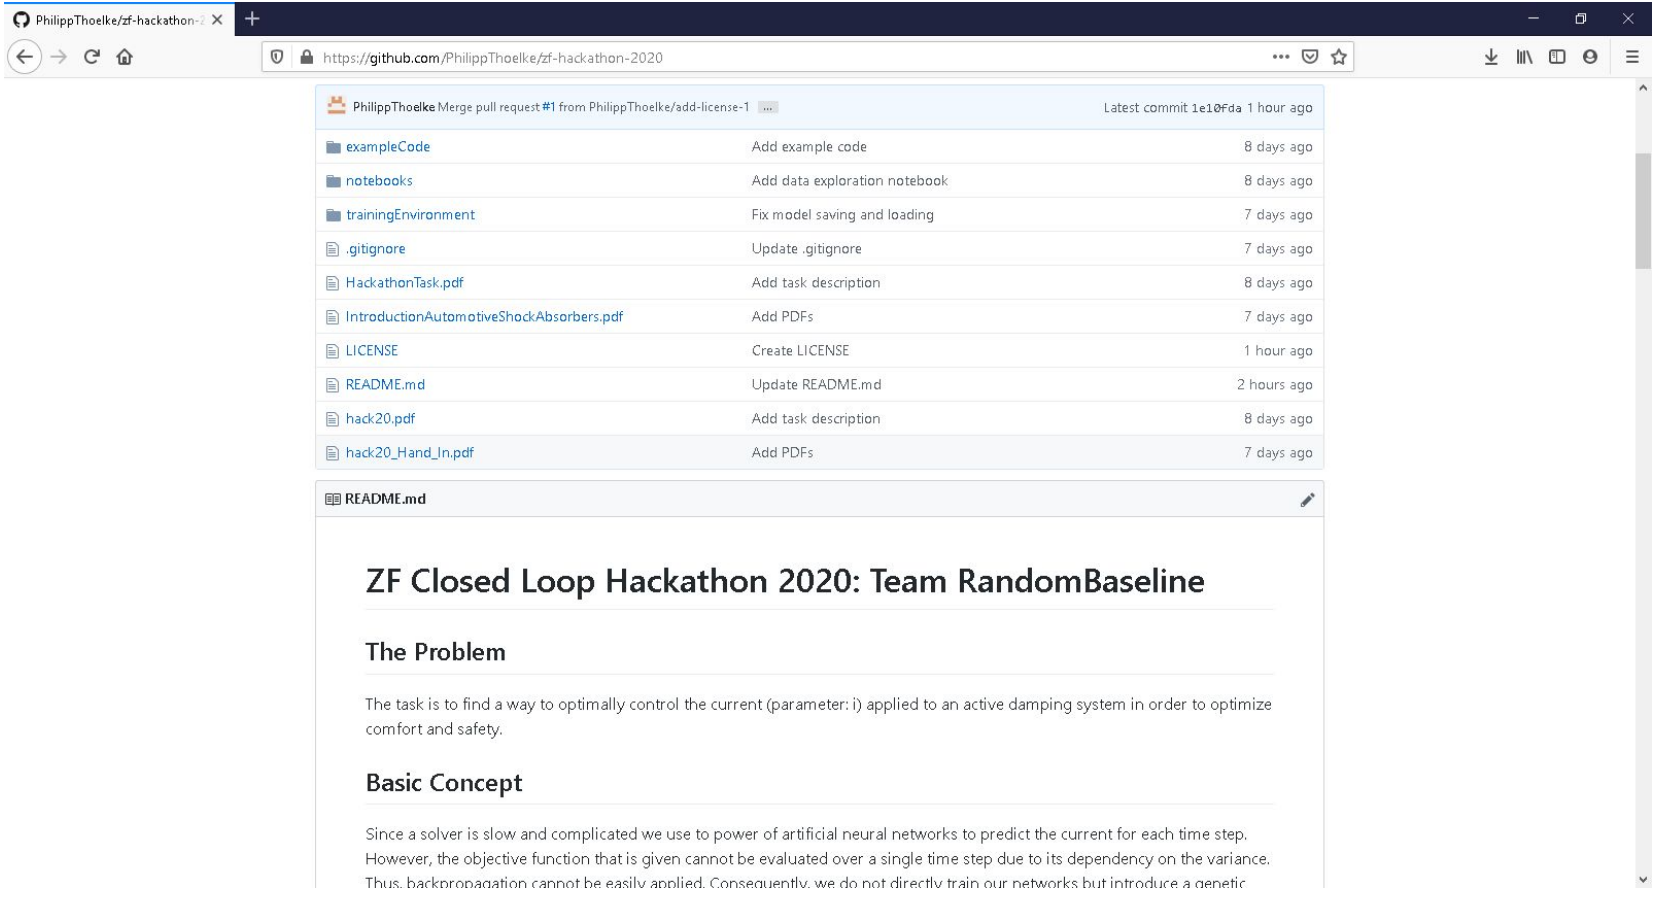

GitHub: https://github.com/PhilippThoelke/zf-hackathon-2020

Thank you.

Questions?

RL continuous action space - Model architecture

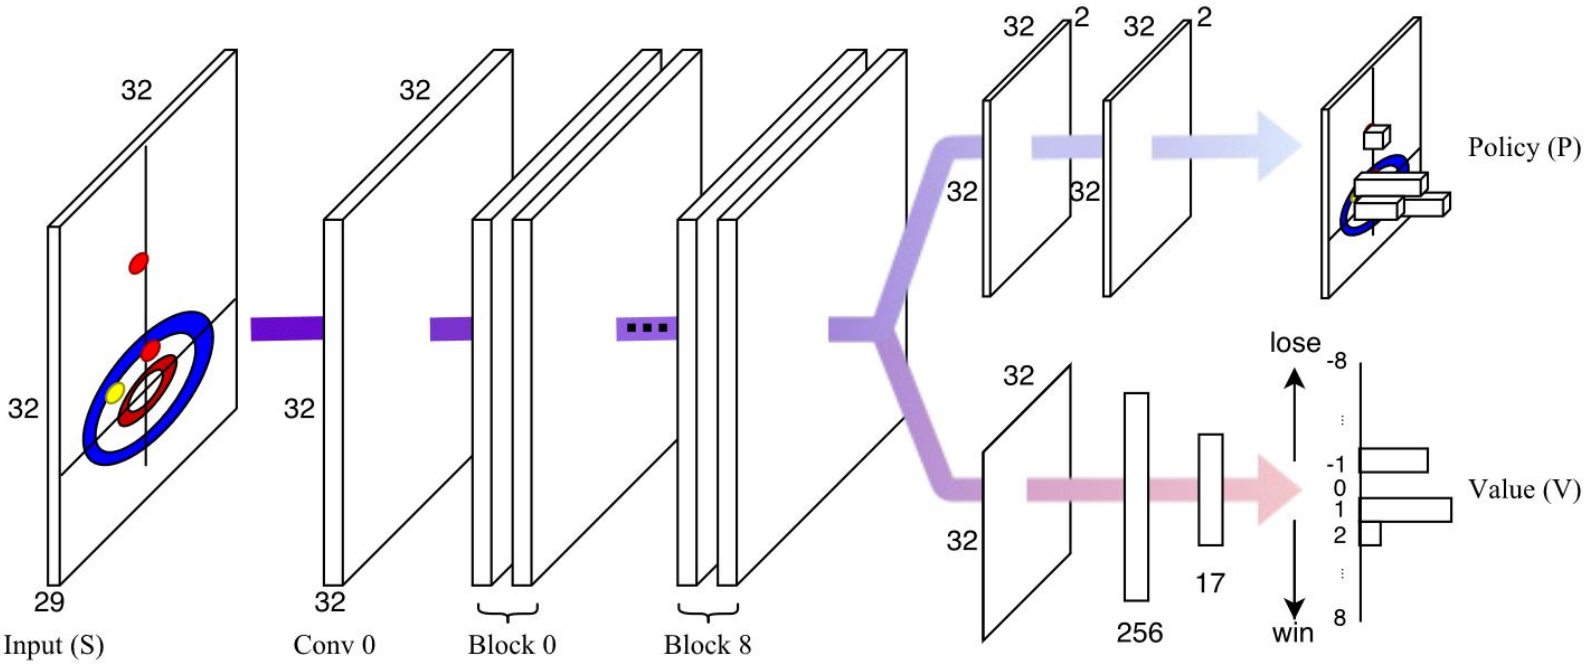

http://proceedings.mlr.press/v80/lee18b/lee18b.pdf